

# Predicting Reasoning Performance Using Ontology Metrics

Yong-Bin Kang<sup>†</sup>, Yuan-Fang Li<sup>†</sup>, Shonali Krishnaswamy<sup>†,§</sup>

<sup>†</sup> Faculty of IT, Monash University, Australia

<sup>§</sup> Institute for Infocomm Research, A\*STAR, Singapore

{yongbin.kang,yuanfang.li,shonali.krishnaswamy}@monash.edu

**Abstract.** A key issue in semantic reasoning is the computational complexity of inference tasks on expressive ontology languages such as OWL DL and OWL 2 DL. Theoretical works have established worst-case complexity results for reasoning tasks for these languages. However, reasoning hardness of individual ontologies has not been adequately characterised. In this paper, we conduct a systematic study to tackle this problem using machine learning techniques, covering over 350 real-world ontologies and four state-of-the-art, widely-used OWL 2 reasoners. Our main contributions are two-fold. Firstly, we learn various classifiers that accurately predict classification time for an ontology based on its metric values. Secondly, we identify a number of ontology metrics that can be used to effectively predict reasoning performance. Our predictive models have been shown to be highly effective, achieving an accuracy of over 80%.

## 1 Introduction

Ontologies are essential building blocks of the Semantic Web. However, the high expressivity of ontology languages also incurs high computational complexity. For example, it has been shown that  $\mathcal{SHOIN}(\mathbf{D})$ , the description logic (DL) underlying OWL DL, is of worst-case NEXPTIME-complete complexity [11]. The complexity of  $\mathcal{SROIQ}(\mathbf{D})$ , the DL underlying OWL 2 DL, is even higher (2NEXPTIME-complete) [7].

The past decade has seen the development of highly optimized inference algorithms for description logics, with tableau algorithms [2] being a leading exemplar. A number of high-performance tableaux-based reasoners have been developed, including FaCT++ [19], HermiT [15], Pellet [16] and TrOWL [18]. Despite the tremendous progress in ontology reasoning, the high theoretical worst-case complexity results for OWL DL and OWL 2 DL still implies that main reasoning problems may be computationally very expensive. For example, it is shown in [4] that although the simpler OWL 2 EL profile has polynomial-time inference algorithms [1], reasoning about large ontologies in OWL 2 EL (Gene Ontology, NCI Thesaurus and SNOMED CT) still requires considerable amount of time and computational resources.

Moreover, worst-case complexity does not necessarily indicate real-world performance on individual ontologies. In this context, it is noteworthy that reasoner benchmarking has been conducted previously [13, 6, 4]. However, these works

only compared inference performance on a small set of ontologies. They did not attempt to draw correlation between characteristics of ontologies with their inference performance. Hence, they do not provide insight into what makes inference difficult on a given ontology.

Metrics have been used widely and successfully to model artefact complexity in combinatorics and software engineering. We believe that it can be similarly applied to the problem of modelling of ontology inference performance, using a set of ontology metrics like those defined in [23] as a basis. In this paper, we tackle the challenge of predicting ontology classification performance by applying machine learning techniques.

Specifically, we conduct a comprehensive and rigorous investigation, using more than 350 real-world ontologies and 4 widely-used OWL 2 DL reasoners (FaCT++, HermiT, Pellet and TrOWL).<sup>1</sup> Multiple classifiers and feature selection algorithms are tested for their effectiveness. Moreover, 27 metrics are studied for their suitability for performance prediction. To the best of our knowledge, to date this is the most comprehensive study on characterizing ontology inference performance, and it is the first study on predicting ontology inference performance.

The main contributions of this paper can be summarized as follows:

**Prediction model** We learn a random forest-based classifier that is consistently accurate in predicting ontology classification time using ontology metrics. The accuracy of the classifier is over 90% for HermiT and TrOWL, and over 80% for FaCT++ and Pellet.

**Key metrics** A set of 8 ontology metrics are identified and ranked to be the most effective in characterizing classification performance. These metrics can provide additional insights into ontology engineering and maintenance.

## 2 Background and Related Work

Works mostly closely related to ours are the ones that compare the performance of OWL reasoners. Benchmarking of description logics (hence ontology) reasoners is not a new topic. However, with the rapid advances made in reasoning algorithms and reasoners, there is sustained interest and need for repeated and rigorous evaluation. Early works [9, 10] mainly used synthetic TBoxes for evaluating system performance on less expressive languages ( $\mathcal{ALC}$  and its predecessors). In [17], Tempich and Volz developed a systematic method of generating synthetic ontologies. They also performed  $k$ -means clustering on 95 medium-sized ontologies (65 class expressions and 25 properties on average) and obtained 3 clusters of ontologies, using a number of different language constructs as features.

The widely-used Lehigh University Benchmark (LUBM) [8] can be used to generate arbitrarily large ABox datasets from a small synthetic university-domain OWL Lite ontology. It has been mainly used to evaluate query answering in triple stores but not OWL reasoners. In [3], 4 ontologies, each from a language

---

<sup>1</sup> Note that TrOWL is a reasoning infrastructure that is capable of performing incomplete reasoning for OWL 2 DL. In our experiment the TrOWL REL reasoner is used to provide OWL 2 DL reasoning through syntactic approximation to OWL 2 EL.

with a different expressivity (RDFS(DL), OWL Lite, OWL DLP and OWL DL) were chosen to compare a number of OWL reasoners. Both TBox and ABox reasoning tasks were compared for a number of reasoners and reason-able triple stores such as Sesame. Reasoner benchmarking has been done using either synthetic or real-world ontologies [13, 6]. More recently, 8 modern reasoners are compared on 3 large OWL 2 EL ontologies (Gene Ontology, NCI Thesaurus and SNOMED CT) [4]. Various dimensions of the OWL reasoners were discussed with performance being a focus. The authors drew the conclusion that there is significant performance variability among reasoners and it should be further investigated. This work partially motivated our investigation in this paper.

In the SEALS project,<sup>2</sup> the Storage and Reasoning Systems Evaluation Campaign 2010 aimed at the evaluation of DL-based reasoners. In the evaluation, the performance of three reasoners FaCT++, HermiT, and jcel were measured and compared in terms of a suite of standard inference services such as classification, class/ontology satisfiability, and logical entailment. Although this evaluation results in a framework revealing a good performance comparison summary of different reasoners, it does not seem to tackle what ontology characteristics impact on the performance of such reasoners.

There has been research on the development of a series of metrics for analyzing ontology complexity. For example, the work [21, 22] proposed some metrics for analyzing ontology complexity by examining the quantity, ratio, and correlativity of classes and relations in a given ontology. However, the metrics developed in this work focused on characteristics of classes without considering a broader range of ontology characteristics. Also, these metrics were mainly designed to analyze complexity evolution and distribution of ontologies, but not for the use of predicting reasoning performance of ontologies. The work [5] defined some metrics to analyze structural complexity of a given ontology. However, it focused only on analyzing coupling between classes as a measure of ontology complexity. Thus, it does not provide any evidence of how the metrics can be used in analyzing reasoning performance of ontologies.

In [23] we proposed a suite of metrics with the aim of characterizing different aspects of ontology *design complexity*. These metrics consider a broader range of ontology characteristics, hence are more suitable for the task of performance prediction. All the metrics can be calculated automatically and efficiently, allowing us to leverage them for predicting reasoning performance.

### 3 Ontology Metrics

In [23] a total of 8 ontology metrics were defined with the aim of measuring different aspects of the design complexity of OWL ontologies. These metrics are defined on a graph representation of an ontology and are used in this paper as a set of features for predicting reasoner performance. They can be divided into two categories: *ontology-level metrics* (ONT) and *class-level metrics* (CLS). In addition to these 8 metrics, we have defined some other metrics that measure

---

<sup>2</sup> <http://www.seals-project.eu>

different aspects of an ontology’s size and structural characteristics. The metrics are defined on the asserted logical axioms in an ontology and they can be divided into two more categories: *anonymous class expressions* (ACE) and *properties* (PRO). For CLS, ACE and PRO metrics, we record the total values of each metric for each ontology. Hence there are 27 distinct metrics in total.

- **Ontology-level Metrics** (ONT). The ONT metrics measure the overall characteristics of a given ontology. Besides the 4 metrics defined previously, including *SOV* (size of vocabulary), *ENR* (edge-node ratio), *TIP* (tree impurity) and *EOG* (entropy of graph) in [23], we define 2 new ONT metrics:
  - *CYC* (Cyclomatic complexity). *CYC* is defined as  $CYC = \#E - \#N + 2 * cc$ , where *cc* is the number of strongly connected components of the ontology represented as a graph. *CYC* measures the number of linearly independent paths in the ontology graph.
  - *RCH* (Expression richness). *RCH* measures the ratio between the number of anonymous class expressions and the total number of class expressions (including named classes).
- **Class-level Metrics** (CLS). Classes are first-class citizens in OWL ontologies, hence we use 4 CLS defined in [23] to capture characteristics of classes in an ontology. These metrics include *NOC* (number of children), *DIT* (depth of inheritance), *CID* (class in-degree), and *COD* (class out-degree).
- **Anonymous Class Expressions** (ACE). ACE are an important ingredient in building up expressive classes. The ACE metrics records, for each kind of anonymous class expression, the number of occurrences in an ontology. There are altogether 9 metrics: enumeration (*ENUM*), negation (*NEG*), conjunction (*CONJ*), disjunction (*DISJ*), universal/existential quantification (*UF/EF*) and min/max/exact cardinality (*MNCAR/MXCAR/CAR*).
- **Properties** (PRO). Similarly, property declarations and axioms may impact reasoning performance. The 8 PRO metrics records the number of occurrences of property declarations and axioms. There are 8 metrics, one each for: object/datatype property declaration (*OBP/DTP*), functional (*FUN*), symmetric (*SYM*), transitive (*TRN*), inverse functional (*IFUN*), property equivalence (*EQV*) and inverse (*INV*).

Note that although there is no metric specifically for ontology expressivity, such information is implicitly captured by the ACE and PRO metrics as 0 for a metric indicates the absence of a particular language construct.

## 4 Investigation Methodology

The principal aims of this paper are two-fold: (1) identifying predictive models that accurately estimate reasoning performance of unknown ontologies, and (2) experimentally discovering significant metrics that influence reasoning performance. The key steps in our investigation can be summarized as follows:

**Scoping.** There are a number of main reasoning tasks on ontologies, including *classification* and *consistency checking*, which are equivalent to each other [2]. In this paper, we focus on classification, the process of making all class subsumption relations explicit in an ontology. Such characteristics make prediction less useful

and reliable. We found that reasoning time for classification is significantly longer than consistency checking, and that there is a significant discrepancy between consistency checking time reported by the reasoners. Thus, we focus on the more difficult reasoning task, classification, and aim to provide insight into predictive models and key metrics embedded in the models. We perform classification on a number of ontologies using different publicly available reasoners. As stated previously, our analysis is conducted on 4 actively-maintained, open-source and widely-used OWL 2 DL reasoners: FaCT++, HermiT, Pellet and TrOWL.

**Data Collection.** We collect a number of ontologies with varying characteristics, which may include the application domain, file size, underlying ontology language, etc. We compute, for each ontology in the collection, (1) its metric values as presented in Section 3, and (2) an average performance time for the reasoning task of *ontology classification* for each of the 4 reasoners.

Furthermore, since our goal is to learn classifiers, the continuous reasoning time values need to be *discretized* in order to assign (i.e. classify) ontologies into separate groups (i.e. class labels) based on their reasoning time.

**Feature Selection.** We hypothesize that different metrics may have different effect on ontology classification performance. *Feature selection* is a very widely-used techniques in data pre-processing to remove irrelevant features. A number of feature selection algorithms are applied to identify and quantitatively study the ontology metrics that have a strong impact on performance. These algorithms typically fall into two categories. *Feature ranking* algorithms rank the features by a metric and eliminate all features that do not achieve an adequate threshold. *Subset selection* algorithms search the set of possible features for the optimal subset. In this work, we consider 6 different feature ranking algorithms (simply features selectors), since we are interested in ranking individual features (metrics) and then finding relevant features based on their ranks. These are the information gain (InfoGain), information gain ratio (GainInfo), support vector machine based weighting (SVM), ReliefF-based (ReliefF), symmetrical uncertainty (Symm), and chi-squared statistic (ChiSquared) feature selectors.

**Predictive Model Development.** In this work, we develop predictive models using classification techniques (in the machine learning sense) to predict reasoning performance of the classification task (in the ontology reasoning sense). In our evaluation, the categories of ontologies are obtained from discretization of reasoning time of the ontologies for the task, as stated above. Each ontology is represented as a pair consisting of a subset of metrics and the corresponding category. The subset of metrics is found using the feature selectors described above. Given a dataset consisting of a set of ontologies, we choose the training and test data based on the standard *10-fold cross validation*, in which each dataset is divided into 10 subsets. Of the 10 subsets, 1 subset is retained as testing data, and the remaining 9 subsets are used as training data. The validation process is then repeated 10 folds (times).

It is well-known that different classifiers tend to produce different prediction performance. Hence, we employ various classifiers and identify the most effective one to build a predictive model for a given dataset. The effectiveness of each

classifier is determined through its *classification accuracy* (simply accuracy), often considered to be the best performance indicator for evaluating classifiers.<sup>3</sup> It measures the proportion of correctly classified ontologies against all ontologies in testing data.

In total, we use 9 representative classifiers that are available in Weka [20], with the aim of finding the best predictive models for the four reasoners. These are classified into 5 categories: Bayesian classifiers (BayesNet (BN) and Naive-Bayes (NB)), decision tree-based classifiers (J48 and RandomForest (RF), REP Tree (RT)), rule-based classifiers (DecisionTable (DT)), a regression-based classifier (SimpleLogistic (SL)), and lazy classifiers (IBk<sub>[1≤k≤10]</sub> and K\*).

**Key Metrics Determination.** Identifying the set of metrics that most highly impact on reasoning time can provide insights for ontology engineering. In this step, such sets of metrics are identified by further analyzing outcomes of the suite of feature selectors utilized in the previous step. More specifically, by examining the metrics used in the classifier chosen in the predictive model for each dataset, we can identify which metrics, in conjunction with a certain classifier, contribute best to leading to the best accuracy.

Given a dataset for each reasoner, we apply the 9 classifiers on various subsets of ontology metrics in the dataset. Recall that such subsets are found using the 6 feature selectors. Then, we identify the best predictive model for the reasoner, which consists of the following three dimensions: (1) a particular classifier leading to the best accuracy, (2) a particular metric subset, used for the classifier, and a specific feature selector that has found the subset, and (3) a prediction performance (accuracy) result carried out by the classifier with the metric subset. We note that the discovered metric subset for each reasoner is regarded as the key metrics leading to its best predictive model. Furthermore, we measure the impact of individual metrics with respect to constructing predictive models for the 4 reasoners based on statistical analysis.

## 5 Data Collection

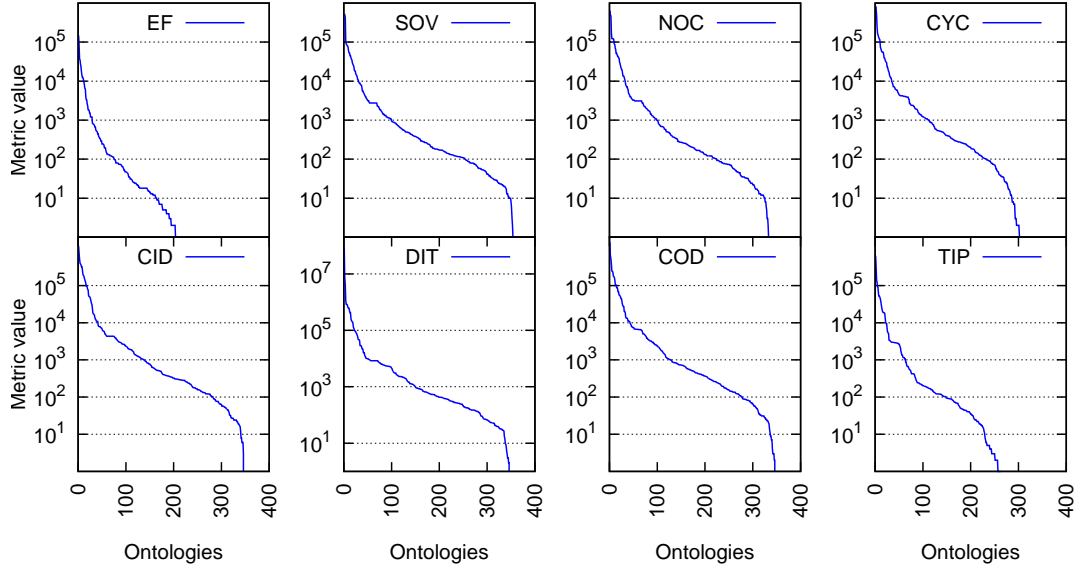
A total of 358 real-world, public-domain ontologies are collected for this work. No preprocessing (cleasing) is done. A large number of these ontologies are collected from the Tones Ontology Repository and NCBO BioPortal.<sup>4</sup> These ontologies vary in file size, ranging from less than 4KB to almost 300MB. However, it is worth noting that file size is not a very good indicator of reasoning performance as a small ontology (such as the DOLCE ontology) may `owl:imports` a large number of other ontologies, which make up the potentially very large import closure that a reasoner considers. Note that all ontologies collected from BioPortal are large, with at least 10,000 *terms*. The expressivity of these ontologies spans OWL 2 EL, OWL DL, through OWL 2 DL to OWL Full. At the same time, this collection also includes some well-known hard ontologies such as FMA, DOLCE, Galen, NCI Thesaurus and the Cell Cycle Ontology (CCO).

<sup>3</sup> F-measure is measured and found to be completely positively correlated to accuracy. For brevity reasons, we only report our experimental results in accuracy.

<sup>4</sup> <http://owl.cs.manchester.ac.uk/repository/>, <http://www.bioontology.org/>

The values of all metrics are calculated; and the distribution of 8 representative metrics are shown in Figure 1, where the metric values are plotted in log scale and ranked by the values. As can be seen quite clearly, the values for these metrics span a large spectrum, ranging from 0 to more than  $10^5$ , and to more than  $10^7$  for *DIT*. Moreover, as expected the majority of ontologies have metric values in the middle of the range, with a few having values closer to the boundary.

Classification time for all ontologies is also collected. All the experiments are performed on a high-performance server running OS Linux 2.6.18 and Java 1.6 on an Intel (R) Xeon X7560 CPU at 2.27GHz with a maximum of 40GB allocated to the 4 reasoner.<sup>5</sup> OWLAPI version 3.2.4 are used to load ontologies and interface with the reasoners. The reasoners that are invoked are: FaCT++ 1.5.3, HermiT 1.3.5, Pellet 2.3.0 and TrOWL 0.8. REL is the underlying reasoner used by TrOWL. These metrics will be revisited in Section 7.

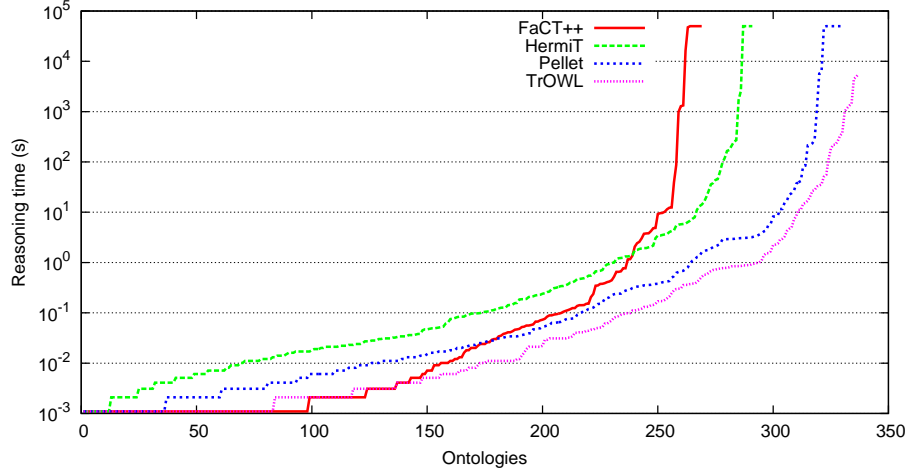


**Fig. 1.** Distributions of values of 8 metrics.

For each ontology and each reasoner, CPU time on ontology classification is averaged over 10 independent runs and recorded. Loading time and pairwise subsumption test time are not included. Trivially simple ontologies (with reasoning time  $\leq 0.01$ s) are later excluded from the experiment to reduce the skewness of the dataset. Some hard ontologies take an extremely long time to classify. Hence, we apply a 50,000-second cutoff for all the reasoners. The distribution of the raw reasoning time for the four reasoners can be found in Figure 2, where classification time is ordered and plotted against the ontologies. It can be observed that FaCT++, HermiT and Pellet all have some ontologies reaching the cut-off time,

<sup>5</sup> To accommodate large ontologies and potential memory leaks in reasoners (due to repeated invocations).

while all ontologies are classified successfully before that for TrOWL. It can also be seen that for relatively easy ontologies ( $\leq 10s$ ), FaCT++ and TrOWL seem to dominate the other 2 reasoners. Compared to performance data reported in [4], the performance on the same ontologies (GO and NCI Thesaurus) seems to be much worse in our experiments, running the same reasoners. Upon closer inspection we notice that the versions of the “same” ontologies are different – we are using more recent versions (current as of November 2011) of these ontologies, which are much large than those versions used in [4].



**Fig. 2.** Raw classification time of the four reasoners.

As stated in the previous section, discretization is a necessary first step before classifiers can be trained. After raw run time values are collected, they are discretized into 5 bins, where the bin ‘T’ contains the trivially simple ontologies (classification time  $\leq 0.01s$ ). The other 4 bins are of unit interval width. The interval width is used as exponent of the reasoning time, i.e.,  $10^i$  is the cutoff point between bin  $i$  and bin  $i + 1$ ,  $1 \leq i \leq 4$ . These 4 bins are labelled ‘A’, ‘B’, ‘C’ and ‘D’. A summary of the discretization and the number of ontologies for each reasoner in each bin is shown in Table 1. It can be seen in the table that each reasoner fails to perform classification on a number of ontologies due to parsing or processing errors or the ontology being inconsistent.

**Table 1.** Discretization of reasoning time and number of ontologies in each bin.

Discretized label	Classification time	FaCT++	HermiT	Pellet	TrOWL
T	$T \leq 0.01s$	161	77	138	188
A	$0.01s < A \leq 1s$	75	154	126	105
B	$1s < B \leq 10s$	16	35	38	17
C	$10s < C \leq 100s$	6	12	12	13
D	$100s < D$	11	13	16	14
<b>Total discretized</b>		269	291	330	337
<b>Ontologies in error</b>		89	67	28	21

It is worth pointing out that the server where the experiments are performed is very well-specced. Although 100 seconds is not a very long time, the same



ontology will take much longer to run on a less powerful computer (mobile devices in particular).

More analysis of the performance characteristics of the reasoners can be found in [12]. All the ontologies, their metric values and reasoning time can be found at [http://www.csse.monash.edu/~yli/metrics\\_perf/](http://www.csse.monash.edu/~yli/metrics_perf/).

## 6 Predictive Models

In this section, we present the first contribution of our work, the construction and analysis of predictive models for classification performance. Our analysis shows that consistently high accuracy ( $> 80\%$ ) is achieved for all the 4 reasoners.

Using 9 classifiers and 6 feature selectors, we learn predictive models as specified in Section 4. For each classifier, the 6 feature selectors are applied to find the best set of metrics. The set of metrics leading to the best accuracy for the classifier and the feature selector is then recorded. The accuracy values of the 9 classifiers are measured. More specifically, an accuracy value is measured for each classifier with 6 different collections of best metrics identified by each of the 6 feature selectors. Eventually, a single set of the best metrics for each (classifier, feature selector) pair is collected.

Section 6.1 presents and analyzes the overall accuracy results of the 4 reasoners. Section 6.2 further characterizes the best predictive model and discusses the effect of feature selection.

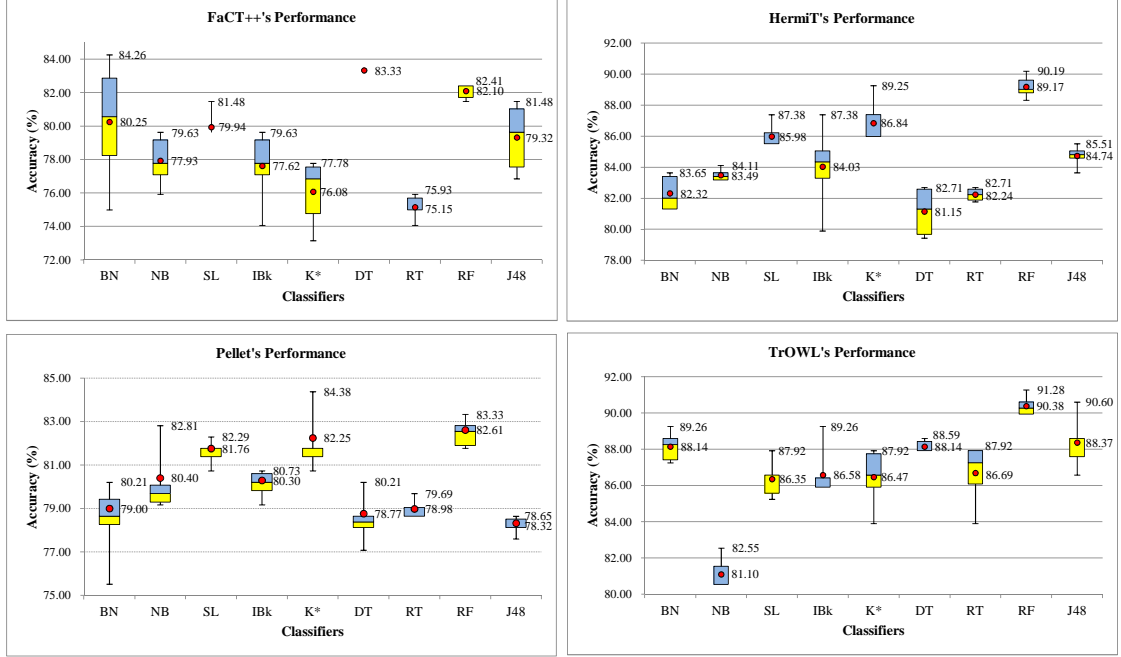
### 6.1 Accuracy Distribution and Analysis

For the 4 reasoners, the accuracy distributions of the 9 classifiers (across the 6 feature selectors) are measured and presented in boxplots in Figure 3. Boxplots provide an excellent visual summary of a distribution through 5 statistical measures: *minimum data value* (MIN), *lower quartile* (Q1), *median* (Q2), *upper quartile* (Q3), *maximum data value* (MAX). Further, we enhance the boxplots by additionally showing the *mean* (AVG) of the accuracy data measured for a single classifier across the 6 feature selectors.

A box itself contains the middle 50% of the accuracy data measured by a classifier with the 6 feature selectors; the upper box area (in blue) denotes the 50th–75th percentile (Q2–Q3) of the data, and the lower box area (in yellow) denotes the 25th–50th percentile (Q1–Q2). The remaining 50% of the data is contained within the areas between the box and the vertical lines or “whiskers”. The ends of the whiskers indicate the minimum and maximum accuracy values. The line inside the box indicates the median value of the accuracy data. The mean of the accuracy is represented by a red circle in the box. Among the above 6 statistical measures, the values of 2 measures, maximum (MAX) and mean (AVG) accuracy, are shown in the plot.

A number of important observations can be made from Figure 3.

- *RF (RandomForest) is the most stable predictive model.* For all the 4 reasoners’ performance, RF has the smallest difference of 1.42 between MIN and MAX, while BN has the largest of 4.57. This indicates that RF leads to the most reliable and stable accuracy results, while BN leads to the most variable accuracy results.



**Fig. 3.** Boxplots displaying the accuracy distribution for the 4 reasoners.

- *Ontology metrics entail good predictive models.* The range of the MAX accuracy for the 4 reasoners is from 84.26 (by BN for FaCT++) to 91.28 (by RF for TrOWL). It indicates that particular subsets of ontology metrics, identified by different feature selectors, can be effectively leveraged for building good predictive models for classification reasoning performance of the reasoners.
- *RF is the best classifier leading to best predictive models for the 4 reasoners.* We examine which of classifiers lead to best predictive models for the 4 reasoners through the statistical analysis of central tendency (e.g. mean) of the measured quantitative values. We compute the mean of the 6 quantitative values shown in Figure 3 across the 4 reasoners. The results are presented in Table 2. The best result in the mean values for each criterion is denoted in **boldface**. The rightmost column, ‘Mean Score’ denotes each classifier’s mean score across the 6 individual scores assigned to the classifier under the 6 values. Table 2 clearly shows that RF leads to the best predictive models for all the 4 reasoners for all the 6 measures. In the following section, we describe these models in more detail.

## 6.2 Best Predictive Models


As each reasoner employs a different set of algorithms and optimization techniques, they may exhibit significantly different performance on the same ontology. As a result, the performance of classifiers may be different for the 4

**Table 2.** The mean quantitative values of the 9 classifiers across the 4 reasoners.

Classifier	6 Statistical Measures						Mean Score
	MIN	Q1	Q2	AVG	Q3	MAX	
BN	81.38	82.89	83.41	83.94	84.74	82.30	83.11
NB	79.85	79.94	80.23	80.79	82.68	80.66	80.69
SL	82.98	83.54	84.12	84.17	85.11	83.47	83.90
IBk	82.54	82.88	83.11	83.51	85.00	82.09	83.19
K*	82.31	83.70	84.16	84.72	86.15	82.83	83.98
DT	82.50	83.04	83.19	83.51	84.40	82.78	83.24
RT	81.27	82.40	82.91	83.43	83.80	80.75	82.43
RF	<b>85.85</b>	<b>85.97</b>	<b>86.41</b>	<b>86.68</b>	<b>87.30</b>	<b>86.03</b>	<b>86.38</b>
J48	82.09	82.90	83.34	83.53	84.63	82.66	83.19

reasoners as well. In this subsection, we further analyze the best classifiers and feature selectors to understand the reasoner-specific behaviours.

As discussed in the previous subsection, RandomForest (RF) is the overall best classifier. This may in part be due to the nature of RF – that it is an *ensemble classifier* that consists of many decision trees. Figure 4 shows, for RF and each reasoner, the MAX classification accuracy (%) for each feature selector, and also their average. The numeric label on top of each bar denotes the number of metrics identified by the corresponding feature selector.



**Fig. 4.** Best predictive models for the 4 reasoners.

RF achieves consistently high accuracy, all higher than 80% for each reasoner with an overall average of 86.03%. For TrOWL, RF achieves 90.38% accuracy on average. It suggests that RF can be effectively used in predicting classification reasoning performance. It also reinforces our belief that ontology metrics can be effective in learning predictive models for the reasoning performance. Moreover, it also opens up the potential to applying our approach to predicting reasoning performance of other reasoners.

It can be observed that each best accuracy result comes with a different number of ontology metrics. The numbers vary from 4 (ReliefF for Pellet) to 25 (ChiSquared for FaCT++). Note that not once is the entire set of 27 metrics

set chosen by any feature selector. This finding establishes the validity of our hypothesis, presented in Section 4, that feature selectors can be leveraged to discover more significant metrics that impact on building more strong predictive models for classification reasoning performance.

## 7 Key Metrics Identification

In this section, we present the second main contribution of this work, the identification of important metrics that have a strong impact on classification performance. Such knowledge can further assist in the task of ontology engineering and maintenance. This identification is achieved through a rigorous quantification of *impact factors* of all 27 ontology metrics used in the classifiers constructed in the previous section.

As discussed in the previous section (Table 2), all the 9 classifiers achieve mean accuracy of at least 80% for all the 4 reasoners. Such high accuracy makes the case for investigating the metrics used by all the classifiers and feature selectors. Two factors has an influence on the significance of a given metric: (1) how frequently it gets selected to be used in the classifiers, and (2) how much it contributes to prediction of reasoning performance. In other words, the more frequently a metric is used in the predictive models (as chosen by the feature selectors), and the more weight it has on the classifiers that use it, the more it influences ontology classification performance. Hence, we combine these two factors to calculate the *impact factor* of all the metrics.

Let metrics be denoted  $m_i, 1 \leq i \leq 27$ , classifiers be denoted by  $c_j, 1 \leq j \leq 9$ , feature selectors be denoted  $f_k, 1 \leq k \leq 6$ , and reasoners be denoted  $r_l, 1 \leq l \leq 4$ . We denote with  $fs_{j,k,l}$  the set of metrics selected for each classifier  $c_j$  by feature selector  $f_k$  for reasoner  $r_l$ . We further denote with  $\#fs_{j,l|i}$  the total number of occurrences of metric  $m_i$  in all metric sets identified by the feature selectors for classifier  $c_j$  and reasoner  $r_l$  ( $0 \leq \#fs_{j,l|i} \leq 6$ ).

Similarly, let  $r_{i,k}^{j,l}$  denote the *weight* of the metric  $m_i$  assigned by feature selector  $f_k$  for the pair  $(c_j, r_l)$  ( $r_{i,k}^{j,l} = 0$  if  $m_i$  is not selected), normalized by  $\max(r_{*,k}^{j,l})$  so that it is between  $[0, 1]$ . We average over all the feature selectors to obtain the average ranked weight  $r_{i,j}^l = \frac{\sum_{k=1}^6 r_{i,k}^{j,l}}{6}$  of  $m_i$ .

Algorithm 1 describes the calculation of the *impact factor* for all the metrics.

For each reasoner (lines 1-8), the combined scores taking into account number of occurrences and weight for each metric are calculated. On line 3, we calculate the total number of occurrences of each metric for each classifier in a  $27 \times 9$  matrix  $mft^l$  (*metric frequency table*). For example, if a metric ‘SOV’ (denoted  $m_2$ ) is in the sets of metrics selected by 4 out of the 6 feature selectors for the classifier RF (denoted  $c_1$ ) and reasoner Pellet (denoted  $r_3$ ), then  $mft_{2,1}^3 = 4$ . Each  $mft_{i,j}^l$  value is then normalized by dividing by the total number of occurrences of all metrics for classifier  $c_j$  on line 4.

From the raw weight  $r(i, j, k, l)$ , we obtain the weight  $r_{i,j}^l$  by averaging over all the 6 feature selectors on line 5. Line 6 then combines the frequency and the averaged weight of each metric by taking the entrywise product of matrices  $mft^l$  and  $r^l$ .

```

Input: Metric number of occurrences  $f(i, j, k, l)$ 
Input: Metric weight  $r(i, j, k, l)$ 
Output: Impact factor for each metric  $\mathbf{mif}_i, 1 \leq i \leq 27$ 
1 foreach reasoner  $r_l$  do
2   Initialize  $27 \times 9$  matrices  $mft^l, mf^l, r^l, nmf^l$ 
3    $mft_{i,j}^l \leftarrow \#fs_{j,l|i}$  /* Metric frequency per classifier */
4    $mf_{i,j}^l \leftarrow \frac{mft_{i,j}^l}{\sum_{i=1}^{27} mft_{i,j}^l}$  /* Normalization */
5    $r_{i,j}^l \leftarrow \frac{\sum_{k=1}^6 r_{i,k}^{j,l}}{6}$  /* Average ranked weight */
6    $mf^l \leftarrow mf^l \circ r^l$  /* Combining frequency with weight */
7    $nmf_{i,j}^l \leftarrow a + (1 - a) \times \frac{mf_{i,j}^l}{\max(mf_{*,j}^l)}$  /* Max frequency normalization */
8 end
9 Initialize  $27 \times 9$  matrix  $\mathbf{nmf}$  for each  $(m_i, c_j)$ ,  $27 \times 1$  vector  $\mathbf{mif}$  for each  $m_i$ 
10  $\mathbf{nmf} \leftarrow \frac{\sum_{l=1}^4 nmf^l}{4}$  /* Average over the reasoners */
11  $\mathbf{mif}_i \leftarrow \frac{\sum_{j=1}^9 \mathbf{nmf}_{i,j}}{9}$  /* Average over the classifiers */
12  $\mathbf{mif}_i \leftarrow \frac{\mathbf{mif}_i}{\max(\mathbf{mif})}$  /* Normalization over max */
13 return  $\mathbf{mif}$ 

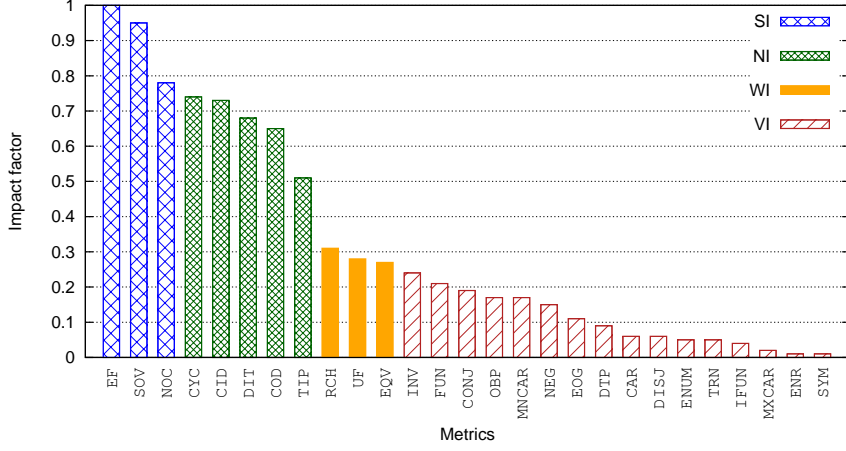
```

**Algorithm 1:** The calculation of the impact factor of metrics.

Note that one problem of the measure  $mf^l$  as calculated on line 6 is that the difference of impact factors between higher and lower frequent metrics tend to be too large. For example, it seems unlikely that 6 occurrences of a metric with a ranked score  $r^l$  in the collection of a classifier in  $mft^l$  truly carry 6 times the significance of a single occurrence with the same or similar ranked score  $r^l$ . To avoid this problem, we apply a normalization technique similar to maximum term frequency normalization [14] to  $mf^l$  on line 7 to obtain the normalized impact factor values for each (metric, classifier) pair for each reasoner. Parameter  $a$  is a value between 0 and 1 and is generally set to 0.4 [14]. It is a *smoothing* term whose role is to dampen the contribution of  $\frac{mf_{i,j}^l}{\max(mf_{*,j}^l)}$ . In this step we also scale down  $mf_{i,j}^l$  by the largest frequency values of all metrics in the collection for a classifier  $c_j$ .

Eventually, the impact factor values are averaged over all 4 reasoners and all 9 classifiers to obtain the final impact factor values of the metrics (lines 10-11).

The metrics can be grouped into a number of categories according to the quartiles their  $\mathbf{mif}$  values fall into: *Strong Impact* (SI,  $0.75 < \mathbf{mif} \leq 1$ ), *Normal Impact* (NI,  $0.5 < \mathbf{mif} \leq 0.75$ ), *Weak Impact* (WI,  $0.25 < \mathbf{mif} \leq 0.5$ ), and *Very Weak Impact* (VI,  $0 \leq \mathbf{mif} \leq 0.25$ ). The ranking and categorization results are shown in Figure 5.



**Fig. 5.** Normalized impact factors of all the ontology metrics.

The 11 metrics in SI, NI and WI are all commonly used in the best predictive models presented in Figure 4. Except one case, all metric sets selected by feature selectors for RF (in Figure 4) are a superset of the 11 metrics in SI, NI and WI in Figure 5. The exception is the set of metrics selected by ReliefF for TrOWL, where there are only 4 metrics. In this case, however, the 4 metrics also belong to these 3 categories. It can be concluded that these 11 metrics form a core group of metrics that are important for predicting classification performance.

Furthermore, it can be observed in Figure 5 that a group of 8 metrics (SI and NI) have high impact on reasoning performance, and that there is a clear separation of **mif** scores between these two groups of metrics and the rest of the metrics. Specifically, it can be clearly seen that (1) the number of existential quantification restrictions (*EF*), (2) the size of an ontology (*SOV*), (3) the number of *independent paths* (*CYC*), (4) the characteristics of named class (the 4 CLS metrics) and (5) the non-treeness of the inheritance graph (*TIP*) have a strong impact on prediction performance.

As mentioned previously, the impact factor indicates a metric’s relative influence on (the prediction of) classification performance. The 8 metrics identified above can hence be used to guide ontology engineering. For example, reducing the size of the ontology (*SOV*), reducing the number of independent paths (*CYC*), reducing the degree of classes (*CID* and *COD*) and making the inheritance graph more tree-like (*TIP*) may significantly improve reasoning performance of the ontology.

## 8 Conclusion

Terminological reasoning has been shown to be a computationally expensive problem, especially for expressive languages such as OWL DL and OWL 2 DL. Despite tremendous progress in the past decade in the design and development of highly optimized algorithms and reasoners, ontology classification is still a very challenging task. It has been shown in previous benchmarking works and in our

own experiments that for certain large and complex ontologies, classification can take an extremely long time (over 50,000 seconds) on a very powerful machine. It is therefore highly desirable to be able to quantitatively analyze and predict reasoning performance using syntactic features.

Metrics have been successfully used to capture different aspects of the syntactic/structural characteristics of various kinds of artefacts (software, combinatorial problems, etc.) including their complexity and empirical hardness. In this paper we propose, develop and evaluate the use of ontology metrics as an effective basis to predict reasoning time for the task of ontology classification. To the best of our knowledge, this is the first known study to apply machine learning techniques (classification) to predict reasoning time for ontologies.

Our contributions in this paper can be summarised as follows: (1) the development of highly effective (over 80% accuracy) predictive models to estimate the reasoning time for an ontology given its metric values, (2) the identification of a set of 8 metrics that have the most impact/correlation with reasoning performance, and (3) a rigorous empirical validation of the proposed methodology with a set of over 350 real-world ontologies. This is the largest study so far in terms of the size of the dataset.

A number of future directions are planned for this work. We will further study the statistical significance of our predictive models and key metrics. Other metrics will be investigated to evaluate their effectiveness in reasoning time prediction. The effect of optimisation techniques on reasoning performance will also be investigated. We also plan to investigate other reasoning tasks (consistency checking) and other machine learning techniques (regression analysis). Lastly, we will study the feasibility of generating synthetic ontologies with specified reasoning performance with metric values as parameters. Such ontologies will be very valuable in the analysis and optimization of reasoning algorithms.

## References

1. F. Baader, S. Brandt, and C. Lutz. Pushing the  $\mathcal{EL}$  envelope further. In K. Clark and P. F. Patel-Schneider, editors, *In Proceedings of the OWLED 2008 DC Workshop on OWL: Experiences and Directions*, 2008.
2. F. Baader and U. Sattler. An overview of tableau algorithms for description logics. *Studia Logica*, 69(1):5–40, 2001.
3. J. Bock, P. Haase, Q. Ji, and R. Volz. Benchmarking OWL reasoners. In *ARea2008 - Workshop on Advancing Reasoning on the Web: Scalability and Commonsense*, June 2008.
4. K. Dentler, R. Cornet, A. ten Teije, and N. de Keizer. Comparison of reasoners for large ontologies in the OWL 2 EL profile. *Semantic Web Journal*, 2(2):71–87, 2011.
5. J. García, F. García, and R. Therón. Defining coupling metrics among classes in an owl ontology. In *Proceedings of the 23rd international conference on Industrial engineering and other applications of applied intelligent systems - Volume Part II*, IEA/AIE’10, pages 12–17, Berlin, Heidelberg, 2010. Springer-Verlag.
6. T. Gardiner, I. Horrocks, and D. Tsarkov. Automated benchmarking of description logic reasoners. In *Proceedings of the 2006 International Workshop on Description Logics (DL2006)*, 2006.

7. B. C. Grau, I. Horrocks, B. Motik, B. Parsia, P. Patel-Schneider, and U. Sattler. OWL 2: The next step for OWL. *Journal of Web Semantics: Science, Services and Agents on the World Wide Web*, 6:309–322, November 2008.
8. Y. Guo, Z. Pan, and J. Heflin. LUBM: A benchmark for OWL knowledge base systems. *Web Semantics: Science, Services and Agents on the World Wide Web*, 3(2-3):158–182, Oct. 2005.
9. J. Heinsohn, D. Kudenko, B. Nebel, and H.-J. Profitlich. An empirical analysis of terminological representation systems. In *Proceedings of the tenth national conference on Artificial intelligence, AAAI’92*, pages 767–773. AAAI Press, 1992.
10. I. Horrocks and P. F. Patel-Schneider. DL systems comparison (summary relation). In *Proceedings of the 1998 International Workshop on Description Logics (DL’98)*, volume 11 of *CEUR Workshop Proceedings*. CEUR-WS.org, 1998.
11. I. Horrocks, P. F. Patel-Schneider, and F. van Harmelen. From *SHIQ* and RDF to OWL: The Making of a Web Ontology Language. *Journal of Web Semantics*, 1(1):7–26, 2003.
12. Y.-B. Kang, Y.-F. Li, and S. Krishnaswamy. A rigorous characterization of reasoning performance – a tale of four reasoners. In *Proceedings of the 1st International Workshop on OWL Reasoner Evaluation (ORE-2012)*, June 2012.
13. Z. Pan. Benchmarking DL reasoners using realistic ontologies. In B. C. Grau, I. Horrocks, B. Parsia, and P. F. Patel-Schneider, editors, *OWLED*, volume 188 of *CEUR Workshop Proceedings*. CEUR-WS.org, 2005.
14. G. Salton and C. Buckley. Term-weighting approaches in automatic text retrieval. *Information Processing and Management: an International Journal*, 24(5):513–523, 1988.
15. R. Shearer, B. Motik, and I. Horrocks. HermiT: A Highly-Efficient OWL Reasoner. In *Proceedings of the 5th International Workshop on OWL: Experiences and Directions (OWLED 2008)*, 2008.
16. E. Sirin, B. Parsia, B. Grau, A. Kalyanpur, and Y. Katz. Pellet: A practical OWL-DL reasoner. *Web Semantics: Science, Services and Agents on the World Wide Web*, 5(2):51–53, June 2007.
17. C. Tempich and R. Volz. Towards a benchmark for semantic web reasoners - an analysis of the DAML ontology library. In Y. Sure and Ó. Corcho, editors, *EON*, volume 87 of *CEUR Workshop Proceedings*. CEUR-WS.org, 2003.
18. E. Thomas, J. Z. Pan, and Y. Ren. TrOWL: Tractable OWL 2 Reasoning Infrastructure. In *ESWC (2)*, pages 431–435. Springer, 2010.
19. D. Tsarkov and I. Horrocks. FaCT++ description logic reasoner: System description. In *Proc. of the Int. Joint Conf. on Automated Reasoning (IJCAR 2006)*, pages 292–297. Springer, 2006.
20. I. H. Witten and E. Frank. *Data mining: Practical machine learning tools and techniques with Java implementations*. Morgan Kaufmann, San Francisco, 2000.
21. Z. YANG, D. Zhang, and C. YE. Evaluation metrics for ontology complexity and evolution analysis. *E-Business Engineering, IEEE International Conference on*, 0:162–170, 2006.
22. D. Zhang, C. Ye, and Z. Yang. An evaluation method for ontology complexity analysis in ontology evolution. In S. Staab and V. Svatek, editors, *Managing Knowledge in a World of Networks*, volume 4248 of *Lecture Notes in Computer Science*, pages 214–221. Springer Berlin / Heidelberg, 2006.
23. H. Zhang, Y.-F. Li, and H. B. K. Tan. Measuring Design Complexity of Semantic Web Ontologies. *Journal of Systems and Software*, 83(5):803–814, 2010.